

Problems that Occur when Multiple Things Use a Database at the Same Time

... and some suggested solutions and workarounds.

A Quick Caveat.



A Quick Caveat.

The Problem

An example.


An example.




An example.






An example.



Store Room









An example.

An example.

```
p = Post.find(123)    SELECT *  
                      FROM posts  
                      WHERE id = 123;
```

An example.

<code>p = Post.find(123)</code>	<code>SELECT *</code>
	<code>FROM posts</code>
	<code>WHERE id = 123;</code>

<code>p.name = "Wow"</code>	
<code>p.save</code>	<code>UPDATE posts</code>
	<code>SET name = "Wow"</code>
	<code>WHERE id = 123;</code>

An example.

```
SELECT *  
FROM posts  
WHERE id = 123;
```

**Something can
happen in here**



```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```


Two examples at once.

A

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

A

```
SELECT *  
FROM posts  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

A

```
SELECT *  
FROM posts  
WHERE id = 123;  
  
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

A

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Nope"  
WHERE id = 123;
```

A

```
SELECT *  
FROM posts  
WHERE id = 123;  
  
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;  
  
UPDATE posts  
SET name = "Nope"  
WHERE id = 123;
```

Winner: B

```
{
```

```
  id: 123,
```

```
  name: "nope",
```

```
  body: "....."
```

```
}
```

ATOMICITY

Atoms

<code>p = Post.find(123)</code>	<code>SELECT *</code>
	<code>FROM posts</code>
	<code>WHERE id = 123;</code>

<code>p.name = "Wow"</code>	
<code>p.save</code>	<code>UPDATE posts</code>
	<code>SET name = "Wow"</code>
	<code>WHERE id = 123;</code>

Atoms

```
p = Post.find(123)
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
p.name = "Wow"  
p.save
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

Atoms

```
p = Post.find(123)
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
p.name = "Wow"  
p.save
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

Atoms

```
p = Post.find(123)
```

```
p.name = "Wow"  
p.save
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

Making Database Interactions Atomic

Three Ways

Three Ways

- I. UPDATE a column's value based on its *current* value.
 - Get the database to figure out the new value.
Don't assume we know what the value is in advance.

1) Push it to the DB.

1) Push it to the DB.

```
p = Post.find(123)    SELECT *  
                      FROM posts  
                      WHERE id = 123;
```

I) Push it to the DB.

```
p = Post.find(123)
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
p.increment(  
    :views, 1  
)
```

```
UPDATE posts  
SET views =  
    COALESCE(views, 0)  
    + 1  
WHERE id = 123;
```

I) Push it to the DB.

```
Post.increment_counter(  
    :views, 1  
)
```

```
UPDATE posts  
SET views =  
    COALESCE(views, 0)  
    + 1  
WHERE id = 123;
```

1) Push it to the DB.

```
p = Post.find(123)          -- SELECT ...

p.tags << "a new tag"      UPDATE posts
                           SET tags =
                           array_append(
                               tags,
                               "a new tag"
                           )

# Rails 4 w/ Postgres      WHERE id = 123;
# and a DB migration,
# eg.
# t.string(
#   :tags,
#   array: true
# )
```

Three Ways

1. UPDATE a column's value based on its *current* value.
 - Get the database to figure out the new value.
Don't assume we know what the value is in advance.
2. Add conditions to the UPDATE.
 - Only update if our assumptions are true.

2) Add Conditions to UPDATE.

2) Add Conditions to UPDATE.

```
p = Post.find(123)          -- SELECT ...

Post.where(
  id: 123,
  name: p.name,
).update_all(
  name: "Wow",
)

UPDATE posts
SET
  name = "Wow"
WHERE
  id = 123 AND
  name = "Old Name";

# => 1 means it worked
# => 0 means it didn't
```

2) Add Conditions to UPDATE.

```
# Database Migration
add_column :posts,
  :lock_version,
  :integer
```


2) Add Conditions to UPDATE.

Database Migration

```
add_column :posts,  
  :lock_version,  
  :integer
```

Form View

```
<%= form.hidden_field  
  :lock_version %>
```

2) Add Conditions to UPDATE.

```
add_column :posts,  
  :lock_version,  
  :integer
```

```
<%= form.hidden_field :lock_version %>
```

2) Add Conditions to UPDATE.

```
add_column :posts,  
  :lock_version,  
  :integer
```

```
<%= form.hidden_field :lock_version %>
```

Controller

```
pp = post_params  
p = Post.find(pp[:id])  
p.lock_version =  
  pp[:lock_version]  
p.name = pp[:name]  
p.save
```

2) Add Conditions to UPDATE.

```
add_column :posts,  
  :lock_version,  
  :integer
```

```
<%= form.hidden_field :lock_version %>
```

Controller

```
pp = post_params
```

```
p = Post.find(pp[:id])
```

```
p.lock_version =
```

```
  pp[:lock_version]
```

```
p.name = pp[:name]
```

```
p.save
```

```
-- SELECT ...
```

2) Add Conditions to UPDATE.

```
add_column :posts,  
  :lock_version,  
  :integer
```

```
<%= form.hidden_field :lock_version %>
```

Controller

```
pp = post_params
```

```
p = Post.find(pp[:id])    -- SELECT ...
```

```
p.lock_version =
```

```
  pp[:lock_version]
```

```
p.name = pp[:name]
```

```
p.save
```

2) Add Conditions to UPDATE.

```
add_column :posts,  
  :lock_version,  
  :integer
```

```
<%= form.hidden_field :lock_version %>
```

Controller

```
pp = post_params  
p = Post.find(pp[:id])  
p.lock_version =  
  pp[:lock_version]  
p.name = pp[:name]  
p.save
```

```
-- SELECT ...  
UPDATE posts  
SET  
  name = "Wow",  
  lock_version = 2  
WHERE  
  id = 123 AND  
  lock_version = 1;
```

2) Add Conditions to UPDATE.

- Problematic...?
 - The "did it update or not?" counter is too coarse.
 - You get another race condition: does the record still exist? Can't differentiate between Stale and Gone.
 - But you can (in most cases) just ask the DB again to check.

Three Ways

1. UPDATE a column's value based on its *current* value.
 - Get the database to figure out the new value.
Don't assume we know what the value is in advance.
2. Add conditions to the UPDATE.
 - Only update if our assumptions are true.
3. Put everything inside a container.
 - Suddenly the *container* is the atom.

3) Put actions inside a container.

Atoms

```
p = Post.find(123)
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
p.name = "Wow"  
p.save
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

Atoms

```
p = Post.find(123)
```

```
p.name = "Wow"
```

```
p.save
```

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

"Just Add a Transaction"

```
ActiveRecord::Base.transaction do |t|
```

```
  p = Post.find(123) # SELECT ...
```

```
  p.name = "Wow"
```

```
  p.save           # UPDATE ...
```

```
end
```





A

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Nope"  
WHERE id = 123;
```

A



```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Wow"  
WHERE id = 123;
```

B

```
SELECT *  
FROM posts  
WHERE id = 123;
```

```
UPDATE posts  
SET name = "Nope"  
WHERE id = 123;
```



3) Put actions inside a container.

3) Put actions inside a container.

- Locking (rows, tables, advisory)

3) Put actions inside a container.

- Locking (rows, tables, advisory)
- Transaction Isolation Levels

Locks

- Locking a particular row.
- Locking an entire table.
- Arbitrary application-level locks.

Locks: SELECT FOR UPDATE

```
p = Post.lock(true).find(123)  
# SELECT ... FOR UPDATE
```

```
p.name = "Foo"  
p.save  
# UPDATEs, COMMITs
```

Locks: Table Lock

- The nuclear options; blocks access entirely.

```
Post.connection.execute(  
    'LOCK TABLE posts IN ... MODE '  
);
```

Locks: Application Lock

https://github.com/heroku/pg_lock

```
post = Post.new(...)
key = "#{Post}-#{post.author}"
PgLock.new(name: key).lock do
  # Stuff, eg. call post.save
  # to insert the new Post.
end
```

Isolation Levels

- Read Uncommitted
- Read Committed (the default)
- Repeatable Read
- Serializable

(Good reference: <http://www.postgresql.org/docs/9.1/static/transaction-iso.html>)

Isolation Levels

- Read Uncommitted
- Read Committed (the default)
- Repeatable Read
- Serializable

(Good reference: <http://www.postgresql.org/docs/9.1/static/transaction-iso.html>)

Isolation Level: Default

```
ActiveRecord::Base.transaction do |t|
```

```
  p = Post.find(123) # SELECT ...
```

```
  p.name = "Wow"
```

```
  p.save              # UPDATE ...
```

```
end
```

Isolation Level: Maximum

```
Post.transaction(  
  isolation: :serializable  
) do |t|  
  
  p = Post.find(123) # SELECT ...  
  
  p.name = "Wow"  
  p.save          # UPDATE ...  
  
end
```

Summin' up.

- In-database modifications where possible.
- Isolation or Locks where not.
- The more restrictions, the slower it's gonna go.
- More work, but makes the problem visible.



Rob 🍷🍷🍷

@damncabbage

Problem: you're missing an example for your talk. You use an e-commerce example. You now have 26 problems.

↩ Reply ★ Favorite ... More

7:47 PM - 8 Sep 2014

Fin.

Rob Howard
@damncabbage

